

МОСКОВСКИЙ ГОСУДАРСТВЕННЫЙ ТЕХНИЧЕСКИЙ УНИВЕРСИТЕТ
им. Н.Э. Баумана

Кафедра «Систем обработки информации и управления»

Лабораторная работа №2
по дисциплине
«Методы машинного обучения»

ИСПОЛНИТЕЛЬ:

группа ИУ5-
23М

Сукач Е.А.
ФИО

подпись

"__" _____ 2020 г.

ПРЕПОДАВАТЕЛЬ:

Гапанюк Ю. Е.
ФИО

подпись

"__" _____ 2020 г.

Москва – 2020

Первая часть

Unique values of features (for more information please see the link above):

- age : continuous.
- workclass : Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.
- fnlwgt : continuous.
- education : Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.
- education-num : continuous.
- marital-status : Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.
- occupation : Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.
- relationship : Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.
- race : White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.
- sex : Female, Male.
- capital-gain : continuous.
- capital-loss : continuous.
- hours-per-week : continuous.
- native-country : United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.
- salary : >50K,<=50K

```
In [6]: import numpy as np
import pandas as pd
```

```
In [7]: data = pd.read_csv('/Users/elizavetasukach/Desktop/MachineLearning/
data.head()
```

Out [7]:

	age	workclass	fnlwgt	education	education- num	marital- status	occupation	relationship	race
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black

1. How many men and women (sex feature) are represented in this dataset?

```
In [8]: data['sex'].value_counts()
```

```
Out [8]: Male      21790
Female    10771
Name: sex, dtype: int64
```

2. What is the average age (age feature) of women?

```
In [25]: #data.loc[data['sex'] == 'Female', 'age'].mean()

data[data['sex']=='Female']['age'].mean()
```

```
Out [25]: 36.85823043357163
```

3. What is the proportion of German citizens (native-country feature)?

```
In [32]: a = float((data['native-country'] == 'Germany').sum()) / data.shape
print("The proportion of German citizens =", "%.2f" % a, "%")
```

The proportion of German citizens = 0.42 %

4-5. What are mean value and standard deviation of the age of those who receive more than 50K per year (salary* feature) and those who receive less than 50K per year? *

```
In [35]: ages1 = data[data['salary']== '>50K']['age']
#ages1 = data.loc[data['salary'] == '>50K', 'age']
#ages2 = data.loc[data['salary'] == '<=50K', 'age']
ages2 = data[data['salary']== '<=50K']['age']
print("The average age of the rich: {0} +- {1} years, poor - {2} +-
      round(ages1.mean()), round(ages1.std(), 1),
      round(ages2.mean()), round(ages2.std(), 1)))
```

The average age of the rich: 44 +- 10.5 years, poor - 37 +- 14.0 years.

6. Is it true that people who receive more than 50k have at least high school education? (education - Bachelors, Prof-school, Assoc-acdm, Assoc-voc, Masters or Doctorate feature)

```
In [48]: mas = data[data['salary']== '>50K']['education'].unique()
element = ['Some-college', '11th', 'HS-grad', '9th', '7th-8th', '12th', '
flag = 0
for i in element:
    if i in mas:
        flag = 1
if (flag==1):
    print("NO")
else:
    print("YES")
```

NO

7. Display statistics of age for each race (race feature) and each gender. Use `groupby()` and `describe()`. Find the maximum age of men of *Amer-Indian-Eskimo* race.

```
In [19]: for (race, sex), sub_df in data.groupby(['race', 'sex']):
    print("Race: {0}, sex: {1}".format(race, sex))
    print(sub_df['age'].describe())
```

```
Race: Amer-Indian-Eskimo, sex: Female
count    119.000000
mean      37.117647
std       13.114991
min       17.000000
25%       27.000000
50%       36.000000
75%       46.000000
max       80.000000
Name: age, dtype: float64
Race: Amer-Indian-Eskimo, sex: Male
count    192.000000
mean     37.208333
```

```
std      12.049563
min      17.000000
25%      28.000000
50%      35.000000
75%      45.000000
max      82.000000
Name: age, dtype: float64
Race: Asian-Pac-Islander, sex: Female
count    346.000000
mean     35.089595
std      12.300845
min      17.000000
25%      25.000000
50%      33.000000
75%      43.750000
max      75.000000
Name: age, dtype: float64
Race: Asian-Pac-Islander, sex: Male
count    693.000000
mean     39.073593
std      12.883944
min      18.000000
25%      29.000000
50%      37.000000
75%      46.000000
max      90.000000
Name: age, dtype: float64
Race: Black, sex: Female
count    1555.000000
mean     37.854019
std      12.637197
min      17.000000
25%      28.000000
50%      37.000000
75%      46.000000
max      90.000000
Name: age, dtype: float64
Race: Black, sex: Male
count    1569.000000
mean     37.682600
std      12.882612
min      17.000000
25%      27.000000
50%      36.000000
75%      46.000000
max      90.000000
Name: age, dtype: float64
Race: Other, sex: Female
count    109.000000
mean     31.678899
std      11.631599
min      17.000000
25%      23.000000
```

```
50%      29.000000
75%      39.000000
max       74.000000
Name: age, dtype: float64
Race: Other, sex: Male
count    162.000000
mean     34.654321
std      11.355531
min      17.000000
25%      26.000000
50%      32.000000
75%      42.000000
max      77.000000
Name: age, dtype: float64
Race: White, sex: Female
count   8642.000000
mean     36.811618
std      14.329093
min      17.000000
25%      25.000000
50%      35.000000
75%      46.000000
max      90.000000
Name: age, dtype: float64
Race: White, sex: Male
count   19174.000000
mean     39.652498
std      13.436029
min      17.000000
25%      29.000000
50%      38.000000
75%      49.000000
max      90.000000
Name: age, dtype: float64
```

8. Among whom the proportion of those who earn a lot(>50K) is more: among married or single men (*marital-status* feature)? Consider married those who have a *marital-status* starting with *Married* (Married-civ-spouse, Married-spouse-absent or Married-AF-spouse), the rest are considered bachelors.

```
In [87]: data1 = data.loc[(data['sex'] == 'Male') &
                        (data['marital-status'].isin(['Never-married',
                                                    'Separated',
                                                    'Divorced',
                                                    'Widowed'])), 'salary'].value_co

proc1 = data1['>50K'] / data.shape[0] * 100
data2 = data.loc[(data['sex'] == 'Male') &
                (data['marital-status'].str.startswith('Married'))], 'salary'].

proc2 = data2['>50K'] / data.shape[0] * 100

print("Процент женатых людей, зарабатывающих больше 50K = ", round(p
print("Процент холостых людей, зарабатывающих больше 50K = ", round(

Процент женатых людей, зарабатывающих больше 50K = 18.3 %
Процент холостых людей, зарабатывающих больше 50K = 2.1 %
```

```
In [22]: data['marital-status'].value_counts()
```

```
Out[22]: Married-civ-spouse      14976
Never-married      10683
Divorced           4443
Separated          1025
Widowed            993
Married-spouse-absent    418
Married-AF-spouse       23
Name: marital-status, dtype: int64
```

9. What is the maximum number of hours a person works per week (*hours-per-week* feature)? How many people work such a number of hours and what is the percentage of those who earn a lot among them?

```
In [88]: max_load = data['hours-per-week'].max()
print("Max time - {0} hours./week.".format(max_load))

num_workaholics = data[data['hours-per-week'] == max_load].shape[0]
print("Total number of such hard workers {0}".format(num_workaholic

rich_share = float(data[(data['hours-per-week'] == max_load)
                        & (data['salary'] == '>50K')].shape[0]) / num_work
print("Percentage of rich among them {0}%".format(int(100 * rich_sh

Max time - 99 hours./week.
Total number of such hard workers 85
Percentage of rich among them 29%
```

10. Count the average time of work (*hours-per-week*) those who earning a little and a lot (*salary*) for each country (*native-country*).

Simple method:

```
In [24]: for (country, salary), sub_df in data.groupby(['native-country', 'salary']):
          print(country, salary, round(sub_df['hours-per-week'].mean(), 2))

? <=50K 40.16
? >50K 45.55
Cambodia <=50K 41.42
Cambodia >50K 40.0
Canada <=50K 37.91
Canada >50K 45.64
China <=50K 37.38
China >50K 38.9
Columbia <=50K 38.68
Columbia >50K 50.0
Cuba <=50K 37.99
Cuba >50K 42.44
Dominican-Republic <=50K 42.34
Dominican-Republic >50K 47.0
Ecuador <=50K 38.04
Ecuador >50K 48.75
El-Salvador <=50K 36.03
El-Salvador >50K 45.0
England <=50K 40.48
England >50K 44.53
France <=50K 41.06
France >50K 50.75
Germany <=50K 39.14
Germany >50K 44.98
Greece <=50K 41.81
Greece >50K 50.62
Guatemala <=50K 39.36
Guatemala >50K 36.67
Haiti <=50K 36.33
Haiti >50K 42.75
Holand-Netherlands <=50K 40.0
Honduras <=50K 34.33
Honduras >50K 60.0
Hong <=50K 39.14
Hong >50K 45.0
Hungary <=50K 31.3
Hungary >50K 50.0
India <=50K 38.23
India >50K 46.48
Iran <=50K 41.44
Iran >50K 47.5
Ireland <=50K 40.95
Ireland >50K 48.0
Italy <=50K 39.62
Italy >50K 45.4
Jamaica <=50K 38.24
Jamaica >50K 41.1
```



```

Japan <=50K 41.0
Japan >50K 47.96
Laos <=50K 40.38
Laos >50K 40.0
Mexico <=50K 40.0
Mexico >50K 46.58
Nicaragua <=50K 36.09
Nicaragua >50K 37.5
Outlying-US(Guam-USVI-etc) <=50K 41.86
Peru <=50K 35.07
Peru >50K 40.0
Philippines <=50K 38.07
Philippines >50K 43.03
Poland <=50K 38.17
Poland >50K 39.0
Portugal <=50K 41.94
Portugal >50K 41.5
Puerto-Rico <=50K 38.47
Puerto-Rico >50K 39.42
Scotland <=50K 39.44
Scotland >50K 46.67
South <=50K 40.16
South >50K 51.44
Taiwan <=50K 33.77
Taiwan >50K 46.8
Thailand <=50K 42.87
Thailand >50K 58.33
Trinidad&Tobago <=50K 37.06
Trinidad&Tobago >50K 40.0
United-States <=50K 38.8
United-States >50K 45.51
Vietnam <=50K 37.19
Vietnam >50K 39.2
Yugoslavia <=50K 41.6
Yugoslavia >50K 49.5

```

Elegant method:

```
In [25]: pd.crosstab(data['native-country'], data['salary'],
                    values=data['hours-per-week'], aggfunc=np.mean).T
```

```
Out [25]:
```

	native-country	?	Cambodia	Canada	China	Columbia	Cuba	Dominican-Republic	I
salary									
<=50K	40.164760	41.416667	37.914634	37.381818	38.684211	37.985714	42.338235	38	
>50K	45.547945	40.000000	45.641026	38.900000	50.000000	42.440000	47.000000	48	

2 rows × 42 columns

In []:

Вторая часть

In [4]: `pip install pandasql`

In [7]:

```
# Import
import pandas as pd
import pandasql as ps
from datetime import datetime
import seaborn
import matplotlib.pyplot as plt
%config InlineBackend.figure_format = 'svg'
from pylab import rcParams
rcParams['figure.figsize'] = 8, 5
```

In [14]: `user_usage = pd.read_csv("/Users/elizavetasukach/Desktop/MachineLea
user_device = pd.read_csv("/Users/elizavetasukach/Desktop/MachineLe
devices = pd.read_csv("/Users/elizavetasukach/Desktop/MachineLearn
devices.rename(columns={"Retail Branding": "manufacturer"}, inplace`

In [4]: `user_usage.head()`

Out [4]:

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id
0	21.97	4.82	1557.33	22787
1	1710.08	136.88	7267.55	22788
2	1710.08	136.88	7267.55	22789
3	94.46	35.17	519.12	22790
4	71.59	79.26	1557.33	22792

In [5]: `user_device.head()`

Out [5]:

	use_id	user_id	platform	platform_version	device	use_type_id
0	22782	26980	ios	10.2	iPhone7,2	2
1	22783	29628	android	6.0	Nexus 5	3
2	22784	28473	android	5.1	SM-G903F	1
3	22785	15200	ios	10.2	iPhone7,2	3
4	22786	28239	android	6.0	ONE E1003	1

In [6]: `devices.head(10)`

Out [6]:

	manufacturer	Marketing Name	Device	Model
0	NaN	NaN	AD681H	Smartfren Andromax AD681H
1	NaN	NaN	FJL21	FJL21
2	NaN	NaN	T31	Panasonic T31
3	NaN	NaN	hws7721g	MediaPad 7 Youth 2
4	3Q	OC1020A	OC1020A	OC1020A
5	7Eleven	IN265	IN265	IN265
6	A.O.I. ELECTRONICS FACTORY	A.O.I.	TR10CS1_11	TR10CS1
7	AG Mobile	AG BOOST 2	BOOST2	E4010
8	AG Mobile	AG Flair	AG_Flair	Flair
9	AG Mobile	AG Go Tab Access 2	AG_Go_Tab_Access_2	AG_Go_Tab_Access_2

In [7]: `result = pd.merge(user_usage, user_device[['use_id', 'platform', 'device']], on='use_id')`
`result.head()`

Out [7]:

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id	platform	device
0	21.97	4.82	1557.33	22787	android	G1950
1	1710.08	136.88	7267.55	22788	android	SG930
2	1710.08	136.88	7267.55	22789	android	SG930
3	94.46	35.17	519.12	22790	android	D230
4	71.59	79.26	1557.33	22792	android	SG36

Выдать id пользователя и девайс, для которого `outgoing_mins_per_month > 1000`

In [20]: `def example1_pandas(user_usage, user_device):`
 `return pd.merge(user_usage[user_usage.outgoing_mins_per_month >`
 `user_device[['use_id', 'device']],`
 `on='use_id') [['use_id', 'device', 'outgoing_mins_per_month']]`

```
In [21]: example1_pandas(user_usage,user_device)
```

```
Out [21]:
```

	use_id	device	outgoing_mins_per_month
0	22788	SM-G930F	1710.08
1	22789	SM-G930F	1710.08
2	22858	ONEPLUS A3003	1221.85

Выдать группировку по полю monthly_mb

```
In [23]: def example2_pandas(user_usage):
         return user_usage.groupby(['monthly_mb']).size()
```

```
In [24]: example2_pandas(user_usage)
```

```
Out [24]: monthly_mb
0.00      1
11.68     1
33.79     1
74.40     1
212.64    1
..
15573.33  9
16611.55  1
20764.45  2
25955.55  1
31146.67  1
Length: 83, dtype: int64
```

```
In [26]: user_usage.shape
```

```
Out [26]: (240, 4)
```

```
In [39]: def example1_pandasql(user_usage,user_device):
         simple_query = '''
             SELECT
             user_usage.use_id,
             user_device.device,
             user_usage.outgoing_mins_per_month
             FROM user_usage JOIN user_device ON user_usage.use_id = use
             WHERE user_usage.outgoing_mins_per_month >1000
             '''
         return ps.sqldf(simple_query, locals())
```

```
In [40]: example1_pandasql(user_usage, user_device)
```

```
Out [40]:
```

	use_id	device	outgoing_mins_per_month
0	22788	SM-G930F	1710.08
1	22789	SM-G930F	1710.08
2	22858	ONEPLUS A3003	1221.85

```
In [25]: def example2_pandasql(user_usage):
    simple_query = '''
        SELECT
        monthly_mb,
        count(*)
        FROM user_usage
        GROUP BY monthly_mb
        '''

    return ps.sqldf(simple_query, locals())

# user_usage.groupby(['monthly_mb']).size()
```

```
In [26]: example2_pandasql(user_usage)
```

```
Out [26]:
```

	monthly_mb	count(*)
0	0.00	1
1	11.68	1
2	33.79	1
3	74.40	1
4	212.64	1
...
78	15573.33	9
79	16611.55	1
80	20764.45	2
81	25955.55	1
82	31146.67	1

83 rows × 2 columns

```
In [41]: import time

def count_mean_time(func, params, N =5):
    total_time = 0
    for i in range(N):
        time1 = time.time()
        if len(params) == 1:
            tmp_df = func(params[0])
        elif len(params) == 2:
            tmp_df = func(params[0], params[1])
        time2 = time.time()
        total_time += (time2 - time1)
    return total_time/N
```

```
In [42]: ex1_times = []
for count in range(1000, 137000, 1000):
    pandasql_time = count_mean_time(example1_pandasql, [user_usage,
    pandas_time = count_mean_time(example1_pandas, [user_usage, user
    ex1_times.append({'count': count, 'pandasql_time': pandasql_tim
```

```
In [43]: ex1_times_df = pd.DataFrame(ex1_times)
ex1_times_df.columns = ['number of rows in daily_engagements', 'pan
ex1_times_df = ex1_times_df.set_index('number of rows in daily_enga
```

```
In [44]: ax = ex1_times_df.plot(title = 'Example #1 time elapsed (seconds)',
<Figure size 576x360 with 2 Axes>
```

```
In [47]: ex2_times = []
for count in range(1000, 137000, 1000):
    pandasql_time = count_mean_time(example2_pandasql, [user_usage]
    pandas_time = count_mean_time(example2_pandas, [user_usage])
    ex2_times.append({'count': count, 'pandasql_time': pandasql_tim
```

```
In [49]: ex2_times_df = pd.DataFrame(ex2_times)
ex2_times_df.columns = ['number of rows in daily_engagements', 'pan
ex2_times_df = ex2_times_df.set_index('number of rows in daily_enga
```

```
In [50]: ax = ex2_times_df.plot(title = 'Example #2 time elapsed (seconds)',
<Figure size 576x360 with 2 Axes>
```

```
In [ ]:
```